

# **Transfer Learning and Hybrid Machine Learning and Deep Learning**


**Slides prepared by: Dr. Sibarama Panigrahi**

# MOTIVATION





# MOTIVATION






# MOTIVATION

Object Tracking




● ● ● / ●

Pose Estimation



● ● / ● ● ●

Object Detection



● ● / ● ●

Action Recognition



● ● ● / ● ●

Autonomous Navigation



● ● ● / ● ●

3D Reconstruction



● ● ● / ●

Crowd Understanding



● ● / ● ● ●

Urban Scene Understanding



● ● / ● ●

Indoor Scene Understanding



● ● ● / ●

Multi-agent Collaboration



● ● / ● ● ● ●

Human Training



● ● ● ●

Aerial Surveying



● ● ● / ● ●

● Image

● Image Label

● Depth/Multi-View

● User Input

● Video

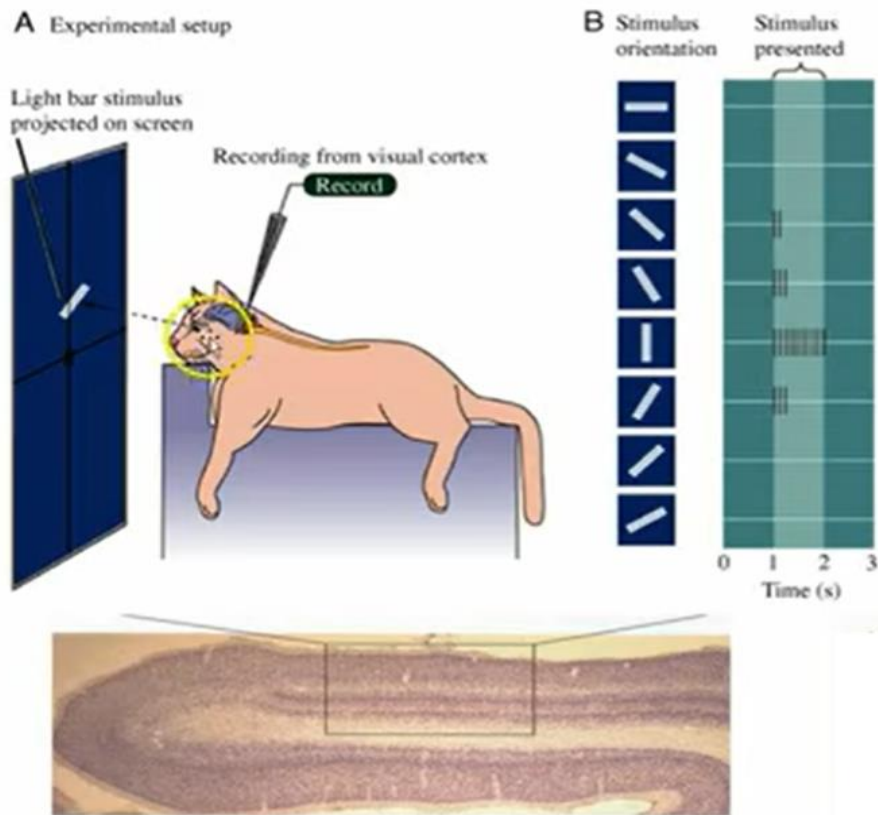
● Physics

● Segmentation/Bounding Box

● Camera Localization

# CONVOLUTIONAL NEURAL NETWORKS

- Popularly called as **CNN or Convnets**.
- [https://en.wikipedia.org/wiki/David\\_H.\\_Hubel#Research](https://en.wikipedia.org/wiki/David_H._Hubel#Research)
- [https://www.youtube.com/watch?v=v20-E\\_2bT2c](https://www.youtube.com/watch?v=v20-E_2bT2c)
- **Hubel and Wiesel** received the Nobel Prize.




# KEY FINDINGS

- Some neurons in the visual cortex fires when lines at specific angle is presented.
- There is a special region called *primary visual cortex that detects edges*.
- There are some more complex neurons that detect motion, depth, color, shapes, complex edges like faces.

## Functional specialization

Match each visual area to its corresponding function:


|     |                               |
|-----|-------------------------------|
| V1  | Motion / <b>Edges</b>         |
| V2  | Stereo                        |
| V3  | Color                         |
| V3a | Texture segregation           |
| V3b | Segmentation, grouping        |
| V4  | Recognition                   |
| V7  | Face recognition              |
| MT  | Attention                     |
| MST | Working memory/mental imagery |
| etc | Etc.                          |





# CONVOLUTION: EDGE DETECTION

We detect edges by applying Convolution operator on the i/p image.



Gray Scale Image  
(6X6)

|     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |

\*

Convolution

Sobel Horizontal  
Edge Detector/  
Kernel

|    |    |    |
|----|----|----|
| 1  | 2  | 1  |
| 0  | 0  | 0  |
| -1 | -2 | -1 |

Kernel / Filter /  
Mask / Operator

# CONVOLUTION: EDGE DETECTION

|     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |

A

$$\begin{aligned}
 &0*1 + 0*2 + 0*1 + \\
 &0*0 + 0*0 + 0*0 + \\
 &-1*0 + -2*0 + -1*0 + \\
 &= 0
 \end{aligned}$$

\*


|    |    |    |
|----|----|----|
| 1  | 2  | 1  |
| 0  | 0  | 0  |
| -1 | -2 | -1 |

K

=

|       |       |       |       |
|-------|-------|-------|-------|
| 0     | 0     | 0     | 0     |
| -1020 | -1020 | -1020 | -1020 |
| -1020 | -1020 | -1020 | -1020 |
| 0     | 0     | 0     | 0     |

R



<http://blog.csdn.net/somTian>



# CONVOLUTION: EDGE DETECTION

|       |       |       |       |
|-------|-------|-------|-------|
| 0     | 0     | 0     | 0     |
| -1020 | -1020 | -1020 | -1020 |
| -1020 | -1020 | -1020 | -1020 |
| 0     | 0     | 0     | 0     |




Normalization

|     |     |     |     |
|-----|-----|-----|-----|
| 255 | 255 | 255 | 255 |
| 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   |
| 255 | 255 | 255 | 255 |



# CONVOLUTION: EDGE DETECTION


- Convolution is a generalization of a dot product.
- We can achieve it through the operation of a neuron in the neural network.



$$V_k = W_{k1} * x_1 + W_{k2} * x_2 + W_{k3} * x_3 + \dots + W_{kn} * x_n$$

# CONVOLUTION: EDGE DETECTION

- Sobel Kernel
  - Horizontal can detect Horizontal Edge
  - Vertical can detect Vertical Edge
- Kernels can be of any size Typically a square matrix.



|   |   |    |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |



# CONVOLUTION: EDGE DETECTION

## Formulation [\[edit\]](#)

The operator uses two 3x3 kernels which are **convolved** with the original image to calculate approximations of the **derivatives** – one for horizontal changes, and one for vertical. If we define **A** as the source image, and **G<sub>x</sub>** and **G<sub>y</sub>** are two images which at each point contain the horizontal and vertical derivative approximations respectively, the computations are as follows:<sup>[2]</sup>

$$\mathbf{G}_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} * \mathbf{A} \quad \text{and} \quad \mathbf{G}_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * \mathbf{A}$$

where \* here denotes the 2-dimensional signal processing **convolution** operation.

Since the Sobel kernels can be decomposed as the products of an averaging and a differentiation kernel, they compute the gradient with smoothing. For example, **G<sub>x</sub>** can be written as

$$\begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} +1 & 0 & -1 \end{bmatrix}$$

The x-coordinate is defined here as increasing in the "right"-direction, and the y-coordinate is defined as increasing in the "down"-direction. At each point in the image, the resulting gradient approximations can be combined to give the gradient magnitude, using:


$$\mathbf{G} = \sqrt{\mathbf{G}_x^2 + \mathbf{G}_y^2}$$

Using this information, we can also calculate the gradient's direction:

$$\angle \mathbf{G}_n \searrow$$



A color picture of a steam engine



The Sobel operator applied to that image

# MOTIVATION FOR PADDING

|     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |

\*

|    |    |    |
|----|----|----|
| 1  | 2  | 1  |
| 0  | 0  | 0  |
| -1 | -2 | -1 |

=

|       |       |       |       |
|-------|-------|-------|-------|
| 0     | 0     | 0     | 0     |
| -1020 | -1020 | -1020 | -1020 |
| -1020 | -1020 | -1020 | -1020 |
| 0     | 0     | 0     | 0     |

- By Performing Convolution operation using 3 X 3 matrix on an input image of 6X6, We have got a 4X4 result.
- By Performing Convolution operation using **K X K** matrix (filter) on an input image of **N X N**, We have got a **N-K+1 X N-K+1** result.
- This results in reduction in Dimension.
- How not to reduce the dimension of the original Image.

# PADDING

|     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |

\*

|    |    |    |
|----|----|----|
| 1  | 2  | 1  |
| 0  | 0  | 0  |
| -1 | -2 | -1 |

=

|       |       |       |       |
|-------|-------|-------|-------|
| 0     | 0     | 0     | 0     |
| -1020 | -1020 | -1020 | -1020 |
| -1020 | -1020 | -1020 | -1020 |
| 0     | 0     | 0     | 0     |

|   |       |       |       |       |   |
|---|-------|-------|-------|-------|---|
| 0 | 0     | 0     | 0     | 0     | 0 |
| 0 | 0     | 0     | 0     | 0     | 0 |
| 0 | -1020 | -1020 | -1020 | -1020 | 0 |
| 0 | -1020 | -1020 | -1020 | -1020 | 0 |
| 0 | 0     | 0     | 0     | 0     | 0 |
| 0 | 0     | 0     | 0     | 0     | 0 |

- This can be achieved by padding 1 row and 1 column of zeros around the resultant matrix. **This is called padding by 1.**



# PADDING

- If you add Zero in padding it is called Zero Padding. (Extensively Used because of Simplicity)
- You can do same value padding.
- Padding  $m$  results in  $n+2m$  size

|   |       |       |       |       |   |
|---|-------|-------|-------|-------|---|
| 0 | 0     | 0     | 0     | 0     | 0 |
| 0 | 0     | 0     | 0     | 0     | 0 |
| 0 | -1020 | -1020 | -1020 | -1020 | 0 |
| 0 | -1020 | -1020 | -1020 | -1020 | 0 |
| 0 | 0     | 0     | 0     | 0     | 0 |
| 0 | 0     | 0     | 0     | 0     | 0 |

|      |       |       |       |       |      |
|------|-------|-------|-------|-------|------|
| 0    | 0     | 0     | 0     | 0     | 0    |
| 0    | 0     | 0     | 0     | 0     | 0    |
| 1020 | -1020 | -1020 | -1020 | -1020 | 1020 |
| 1020 | -1020 | -1020 | -1020 | -1020 | 1020 |
| 0    | 0     | 0     | 0     | 0     | 0    |
| 0    | 0     | 0     | 0     | 0     | 0    |

# STRIDE

|     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |

A

\*

|    |    |    |
|----|----|----|
| 1  | 2  | 1  |
| 0  | 0  | 0  |
| -1 | -2 | -1 |

K

=

|       |       |       |       |
|-------|-------|-------|-------|
| 0     | 0     | 0     | 0     |
| -1020 | -1020 | -1020 | -1020 |
| -1020 | -1020 | -1020 | -1020 |
| 0     | 0     | 0     | 0     |

R

Stride 1= Shift by 1 column/row

# STRIDE

|     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |

A

\*

|    |    |    |
|----|----|----|
| 1  | 2  | 1  |
| 0  | 0  | 0  |
| -1 | -2 | -1 |

K

=

|       |       |       |       |
|-------|-------|-------|-------|
| 0     | 0     | 0     | 0     |
| -1020 | -1020 | -1020 | -1020 |
| -1020 | -1020 | -1020 | -1020 |
| 0     | 0     | 0     | 0     |

R

Stride 2= Shift by 2 column/row



# STRIDE

|     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 0   | 0   | 0   | 0   | 0   | 0   |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |
| 255 | 255 | 255 | 255 | 255 | 255 |

A

\*

|    |    |    |
|----|----|----|
| 1  | 2  | 1  |
| 0  | 0  | 0  |
| -1 | -2 | -1 |

K

=

|       |       |       |       |
|-------|-------|-------|-------|
| 0     | 0     | 0     | 0     |
| -1020 | -1020 | -1020 | -1020 |
| -1020 | -1020 | -1020 | -1020 |
| 0     | 0     | 0     | 0     |


R

 $n \times n$ 


Stride=s and Kernel= kX k

 $\left(\left\lfloor \frac{n-k}{s} \right\rfloor + 1\right) \times \left(\left\lfloor \frac{n-k}{s} \right\rfloor + 1\right)$

# CONVOLUTION IN AN RGB IMAGE



# CONVOLUTION IN AN RGB IMAGE



# CONVOLUTION IN AN RGB IMAGE

|     |     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   | 0   | 0   | ... |
| 0   | 156 | 155 | 156 | 158 | 158 | ... |
| 0   | 153 | 154 | 157 | 159 | 159 | ... |
| 0   | 149 | 151 | 155 | 158 | 159 | ... |
| 0   | 146 | 146 | 149 | 153 | 158 | ... |
| 0   | 145 | 143 | 143 | 148 | 158 | ... |
| ... | ... | ... | ... | ... | ... | ... |

Input Channel #1 (Red)

|     |     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   | 0   | 0   | ... |
| 0   | 167 | 166 | 167 | 169 | 169 | ... |
| 0   | 164 | 165 | 168 | 170 | 170 | ... |
| 0   | 160 | 162 | 166 | 169 | 170 | ... |
| 0   | 156 | 156 | 159 | 163 | 168 | ... |
| 0   | 155 | 153 | 153 | 158 | 168 | ... |
| ... | ... | ... | ... | ... | ... | ... |

Input Channel #2 (Green)

|     |     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   | 0   | 0   | ... |
| 0   | 163 | 162 | 163 | 165 | 165 | ... |
| 0   | 160 | 161 | 164 | 166 | 166 | ... |
| 0   | 156 | 158 | 162 | 165 | 166 | ... |
| 0   | 155 | 155 | 158 | 162 | 167 | ... |
| 0   | 154 | 152 | 152 | 157 | 167 | ... |
| ... | ... | ... | ... | ... | ... | ... |

Input Channel #3 (Blue)

|    |    |    |
|----|----|----|
| -1 | -1 | 1  |
| 0  | 1  | -1 |
| 0  | 1  | 1  |

Kernel Channel #1



158

|   |    |    |
|---|----|----|
| 1 | 0  | 0  |
| 1 | -1 | -1 |
| 1 | 0  | -1 |

Kernel Channel #2



-14

|   |    |   |
|---|----|---|
| 0 | 1  | 1 |
| 0 | 1  | 0 |
| 1 | -1 | 1 |

Kernel Channel #3



653

+

+

+ 1 = 798

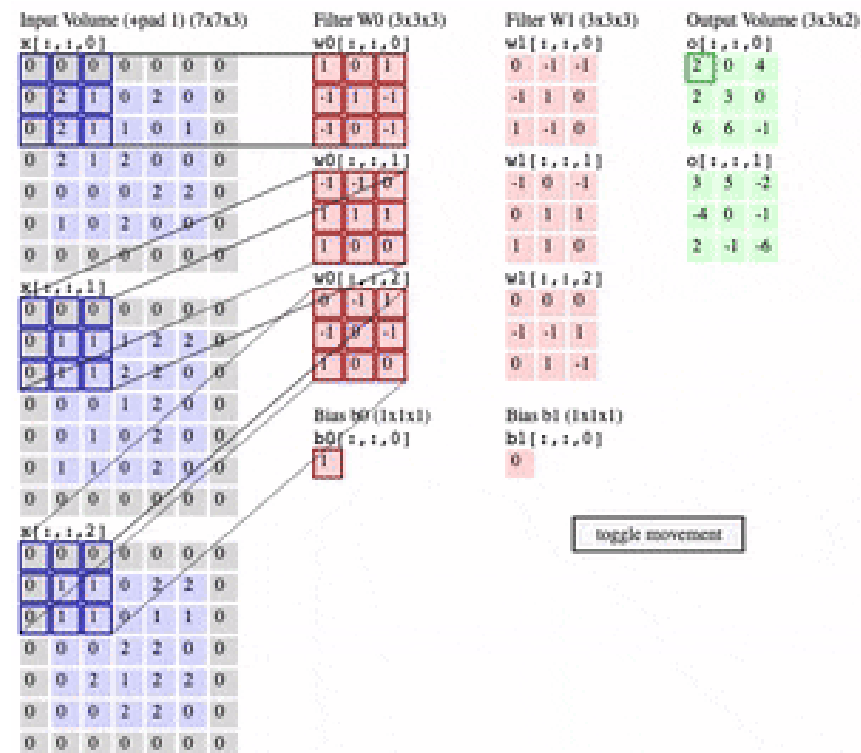
  
Bias = 1

Output

|     |     |     |     |     |
|-----|-----|-----|-----|-----|
| -25 | 466 | 466 | 475 | ... |
| 295 | 787 | 798 |     | ... |
|     |     |     |     | ... |
|     |     |     |     | ... |
| ... | ... | ... | ... | ... |

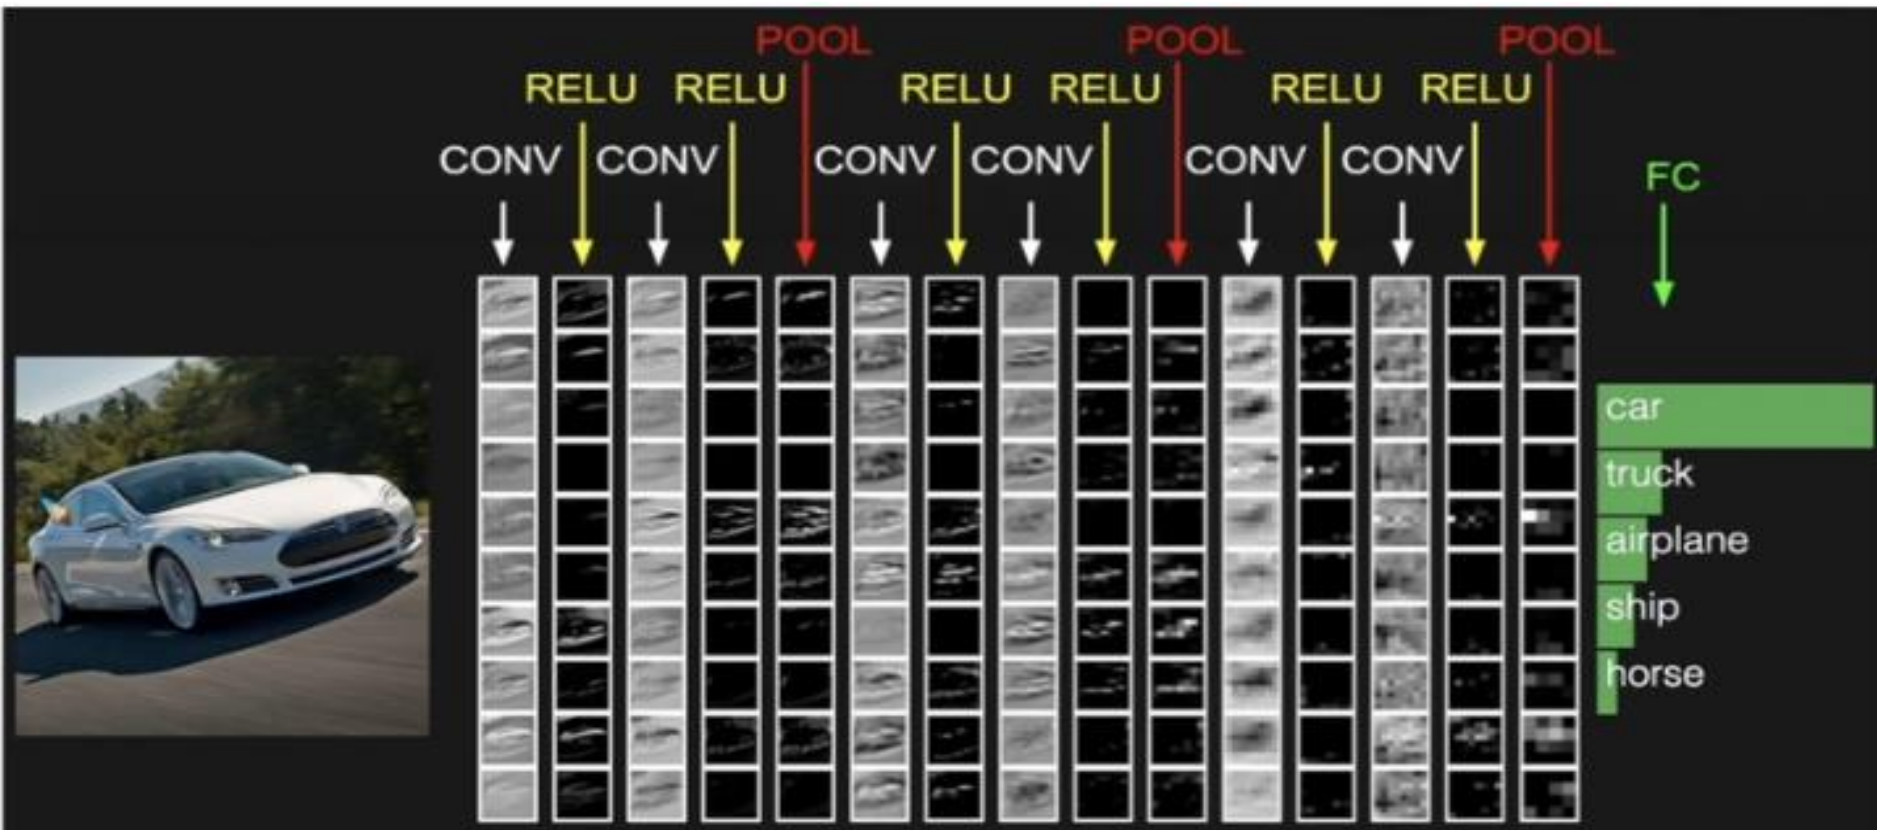
# CONVOLUTION LAYER IN CNN

- Biologically Inspired
- Multiple Edge Detectors: Multiple Kernels
- In MLP, we learn the weights
- In CNN, We learn the kernel matrices



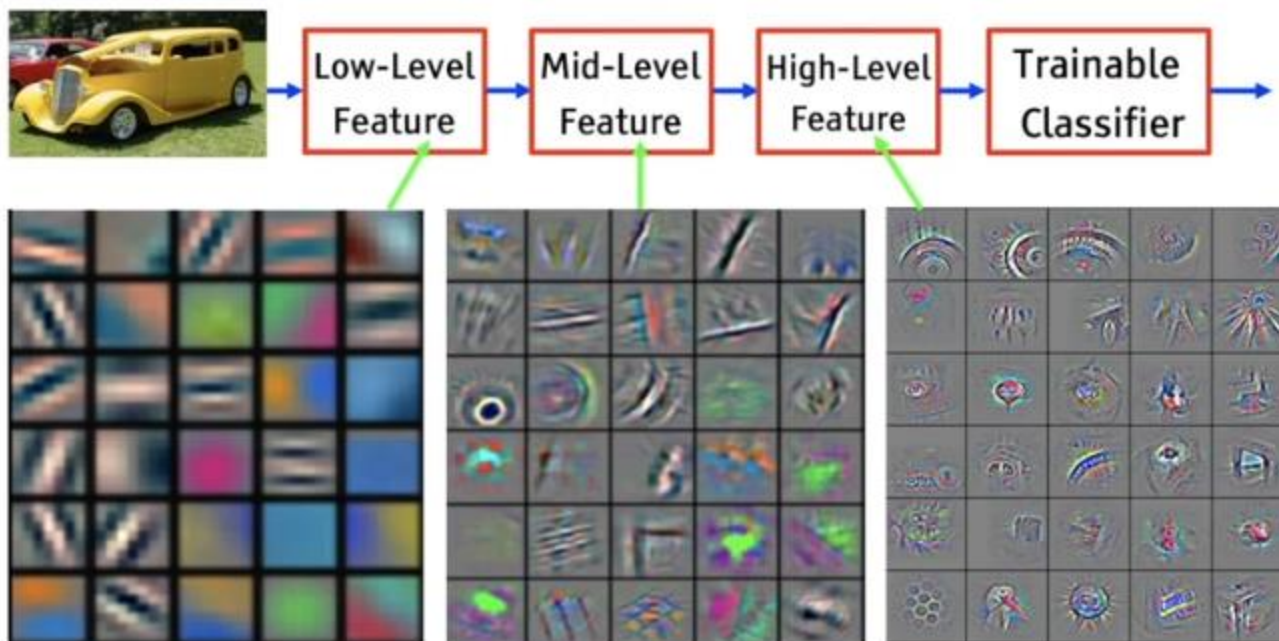


# CONVOLUTION LAYER IN CNN




# WHY MULTIPLE LAYERS?

- Hierarchy of representations with increasing level of abstraction
- Each stage is a kind of trainable feature transform
- **Image recognition:** Pixel → edge → texton → motif → part → object
- **Text:** Character → word → word group → clause → sentence → story
- **Speech:** Sample → spectral band → sound → ... → phone → phoneme → word



# MOTIVATION FOR POOLING

- **Location Invariant:** Changing the location will not change the object.



# MOTIVATION FOR POOLING

- **Scale Invariant:** Subsampling Image will not change the Image

bird






Subsampling

bird



# MOTIVATION FOR POOLING

- Rotation Invariant:** Rotating an image will not change the Object.





# MOTIVATION FOR POOLING


- **Pooling is a concept** that makes the CNN models **invariant to *Location, Scale and Rotation***.

# POOLING

- **E.g. Max Pooling**

- Let's have a 4X4 image with kernel size 2X2 and Stride 2

- Popular method



# HOW DERIVATIVE OF POOLING WORKS IN CNN

## Derivatives of Pooling

Pooling layer subsamples statistics to obtain summary statistics with any aggregate function (or filter)  $g$  whose input is vector, and output is scalar. Subsampling is an operation like convolution, however  $g$  is applied to disjoint (non-overlapping) regions.

### ■ Definition: *subsample (or downsample)*

Let  $m$  be the size of pooling region,  $x$  be the input, and  $y$  be the output of the pooling layer.  $\text{subsample}(f, g)[n]$  denotes the  $n$ -th element of  $\text{subsample}(f, g)$ .

$$y_n = \text{subsample}(x, g)[n] = g(x_{(n-1)m+1:m})$$


$$y = \text{subsample}(x, g) = [y_n]$$

$$g(x) = \begin{cases} \frac{\sum_{k=1}^m x_k}{m}, & \frac{\partial g}{\partial x} = \frac{1}{m} \\ \max(x), & \frac{\partial g}{\partial x_i} = \begin{cases} 1 & \text{if } x_i = \max(x) \\ 0 & \text{otherwise} \end{cases} \\ \|x\|_p = \left( \sum_{k=1}^m |x_k|^p \right)^{1/p}, & \frac{\partial g}{\partial x_i} = \left( \sum_{k=1}^m |x_k|^p \right)^{1/p-1} |x_i|^{p-1} \\ \text{or any other differentiable } \mathbf{R}^m \rightarrow \mathbf{R} \text{ functions} \end{cases}$$

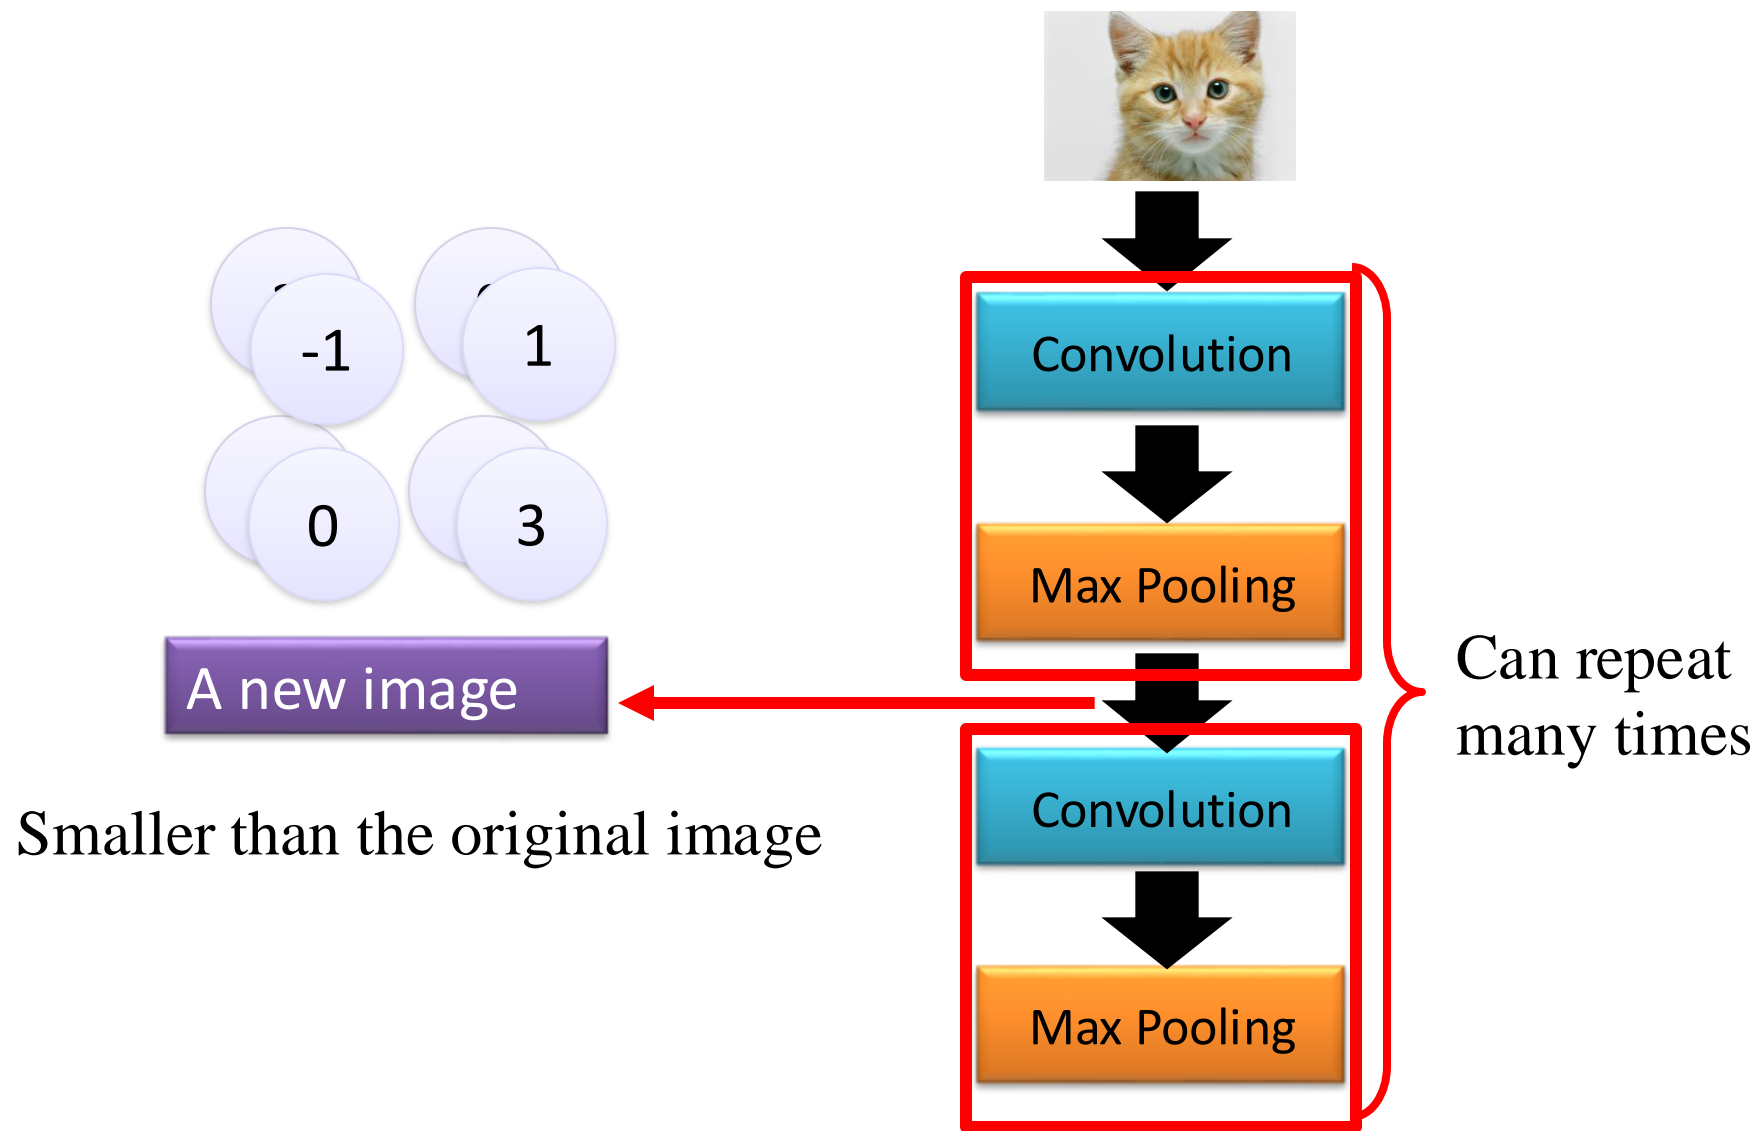
mean pooling

max pooling

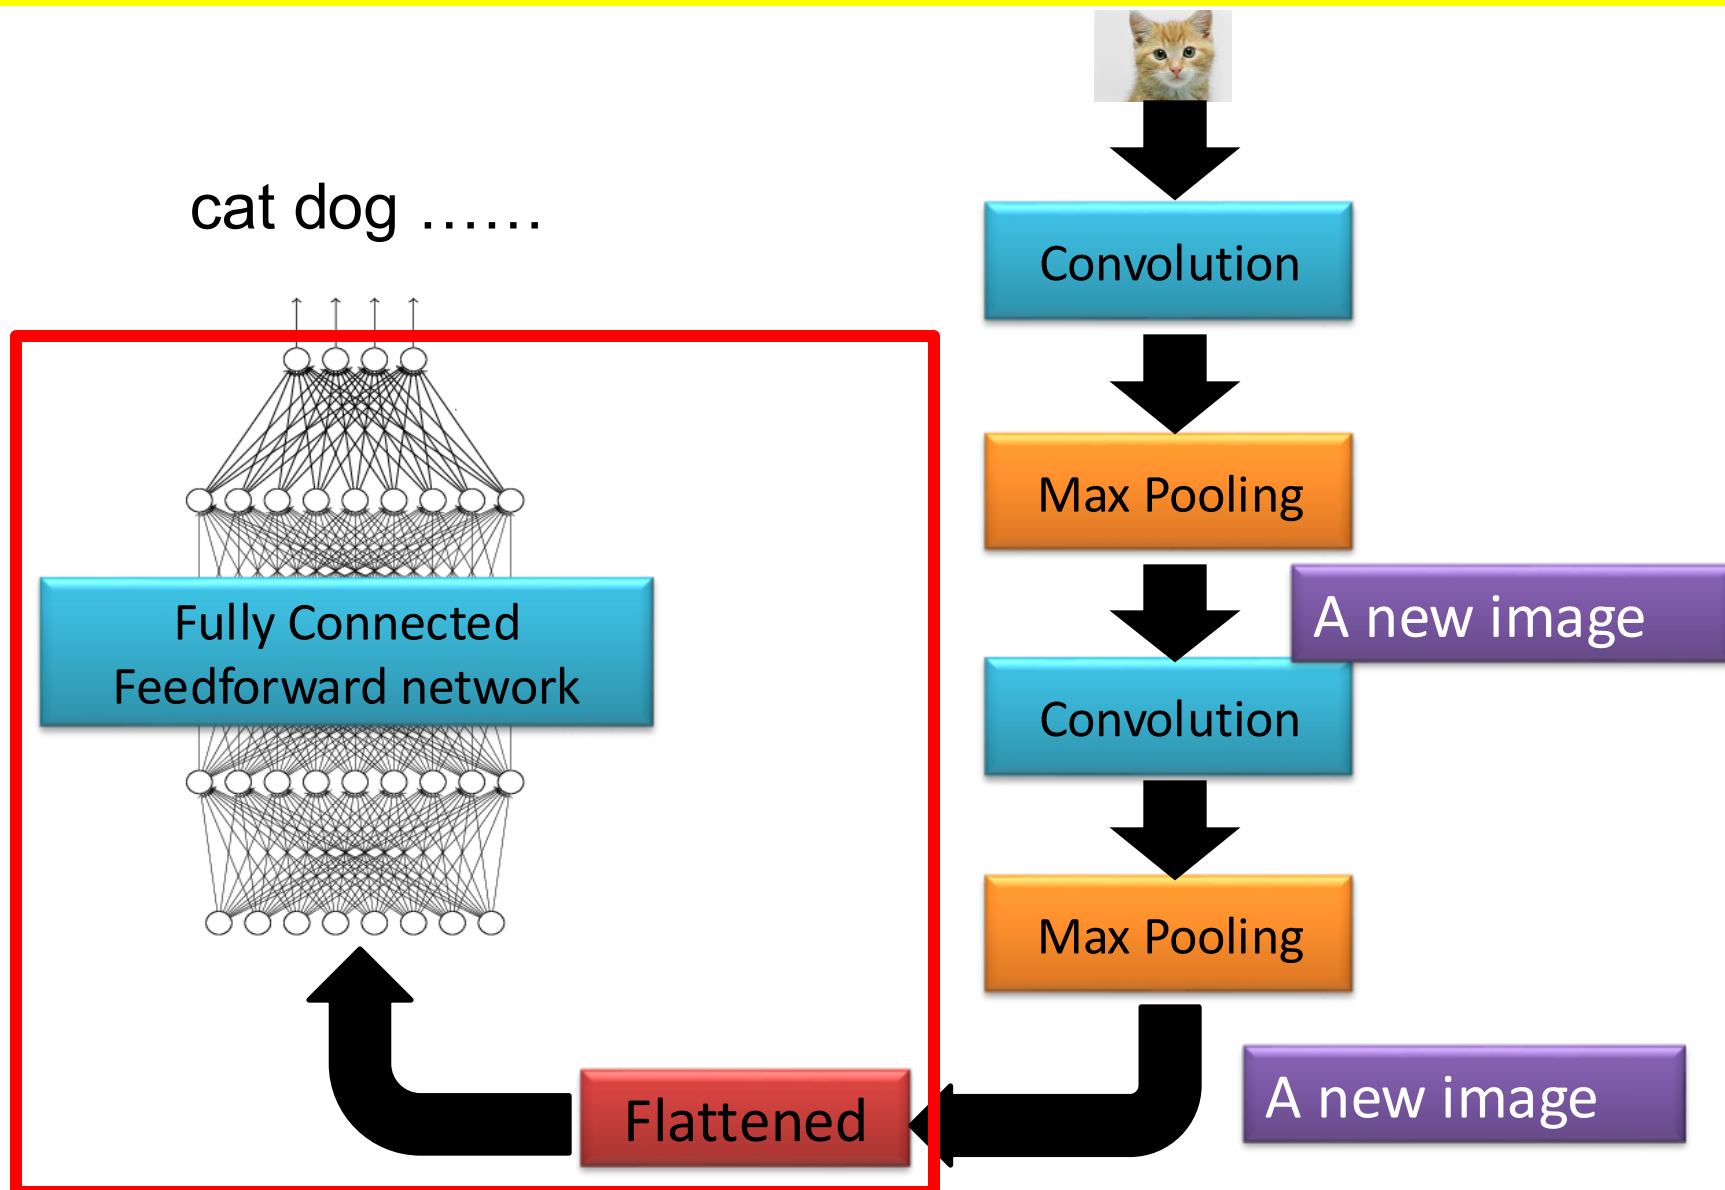
$L^p$  pooling



# CONVNETS



# CONVNETS





# POPULAR CNN MODELS

- LeNet
- AlexNet
- VGGNet
  - VGG16
  - VGG19
- ResNets
- GoogLeNet

# LENET

- [https://d2l.ai/chapter\\_convolutional-neural-networks/lenet.html](https://d2l.ai/chapter_convolutional-neural-networks/lenet.html)




Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Y., Bottou, L., Bengio, Y., & Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278-2324.


# LENET

[https://d2l.ai/chapter\\_convolutional-neural-networks/lenet.html](https://d2l.ai/chapter_convolutional-neural-networks/lenet.html)



Fig. 6.6.2 Compressed notation for LeNet-5.

## ALEXNET

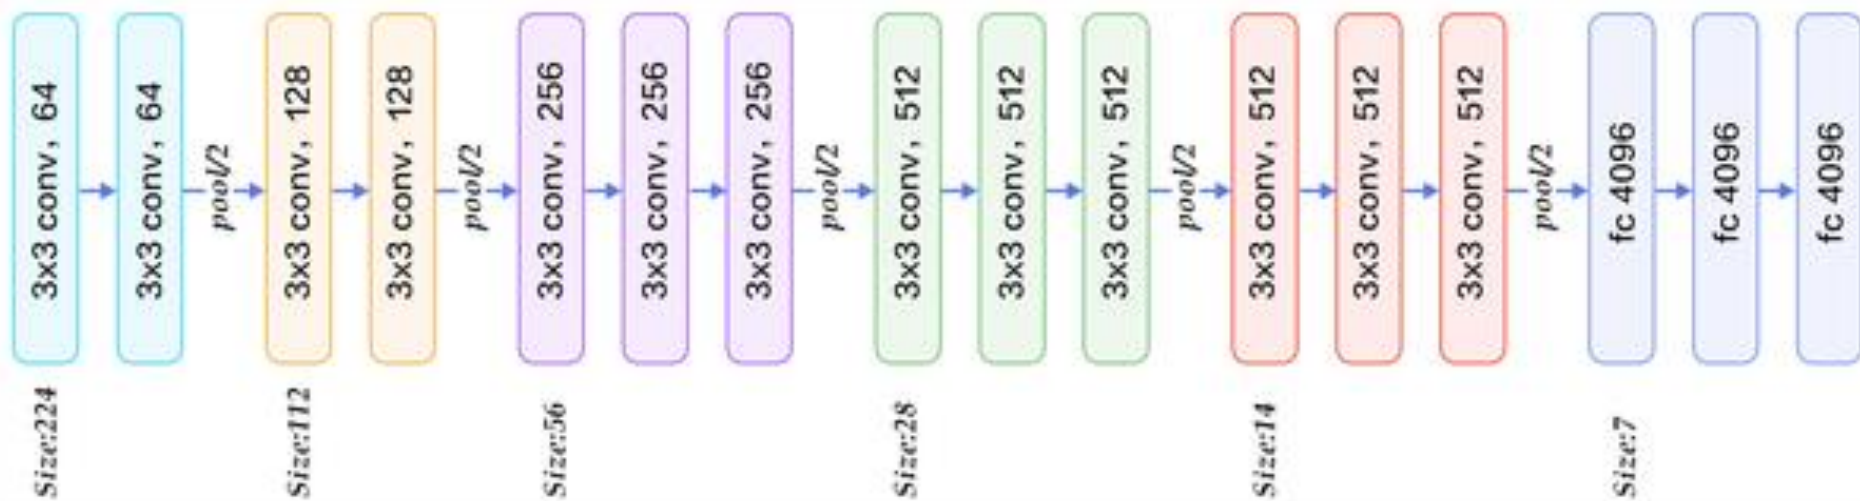


Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). ImageNet Classification with Deep Convolutional Neural Networks. NIPS'12 Proceedings of the 25th International Conference on Neural Information Processing Systems, Lake Tahoe, Nevada, 3-6 December. 1: 1097-1105.

# VGGNET- VGG16

All Convolution (3X3 Kernel, Stride=1, Padding='Same')


All MaxPool (2X2, Stride=2)



Simonyan, K., & Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*.



# VGGNET- VGG16





# VGG-16

- Reference

[https://github.com/keras-team/keras-applications/blob/master/keras\\_applications/vgg16.py](https://github.com/keras-team/keras-applications/blob/master/keras_applications/vgg16.py)

# RESIDUAL NETWORKS: RESNETS



He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 770-778).

# RESIDUAL NETWORKS: RESNETS


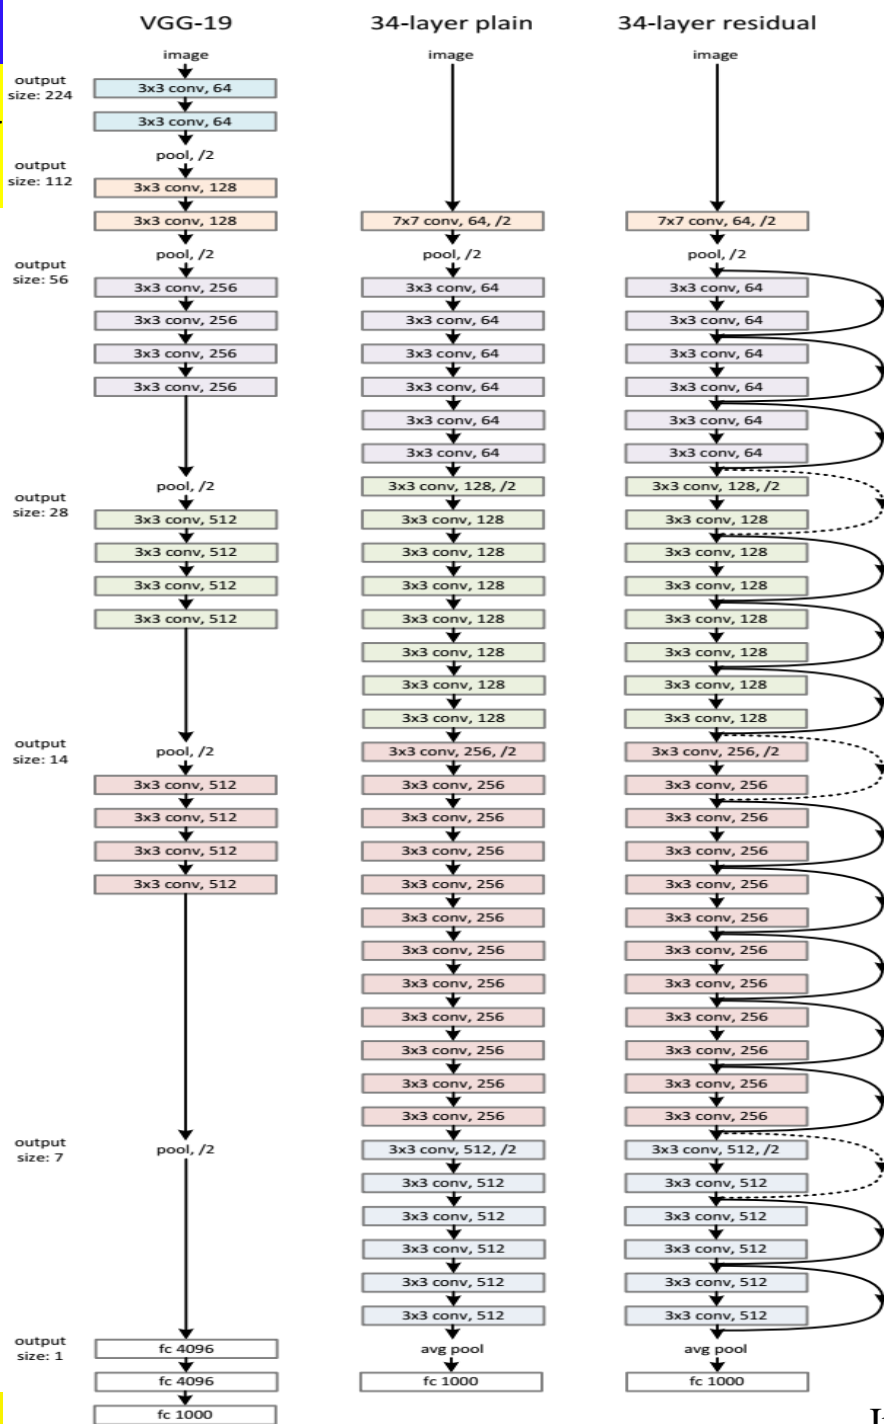




Figure 2. Residual learning: a building block.

He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 770-778).

# RESIDU



# RESIDUAL NETWORKS: RESNETS

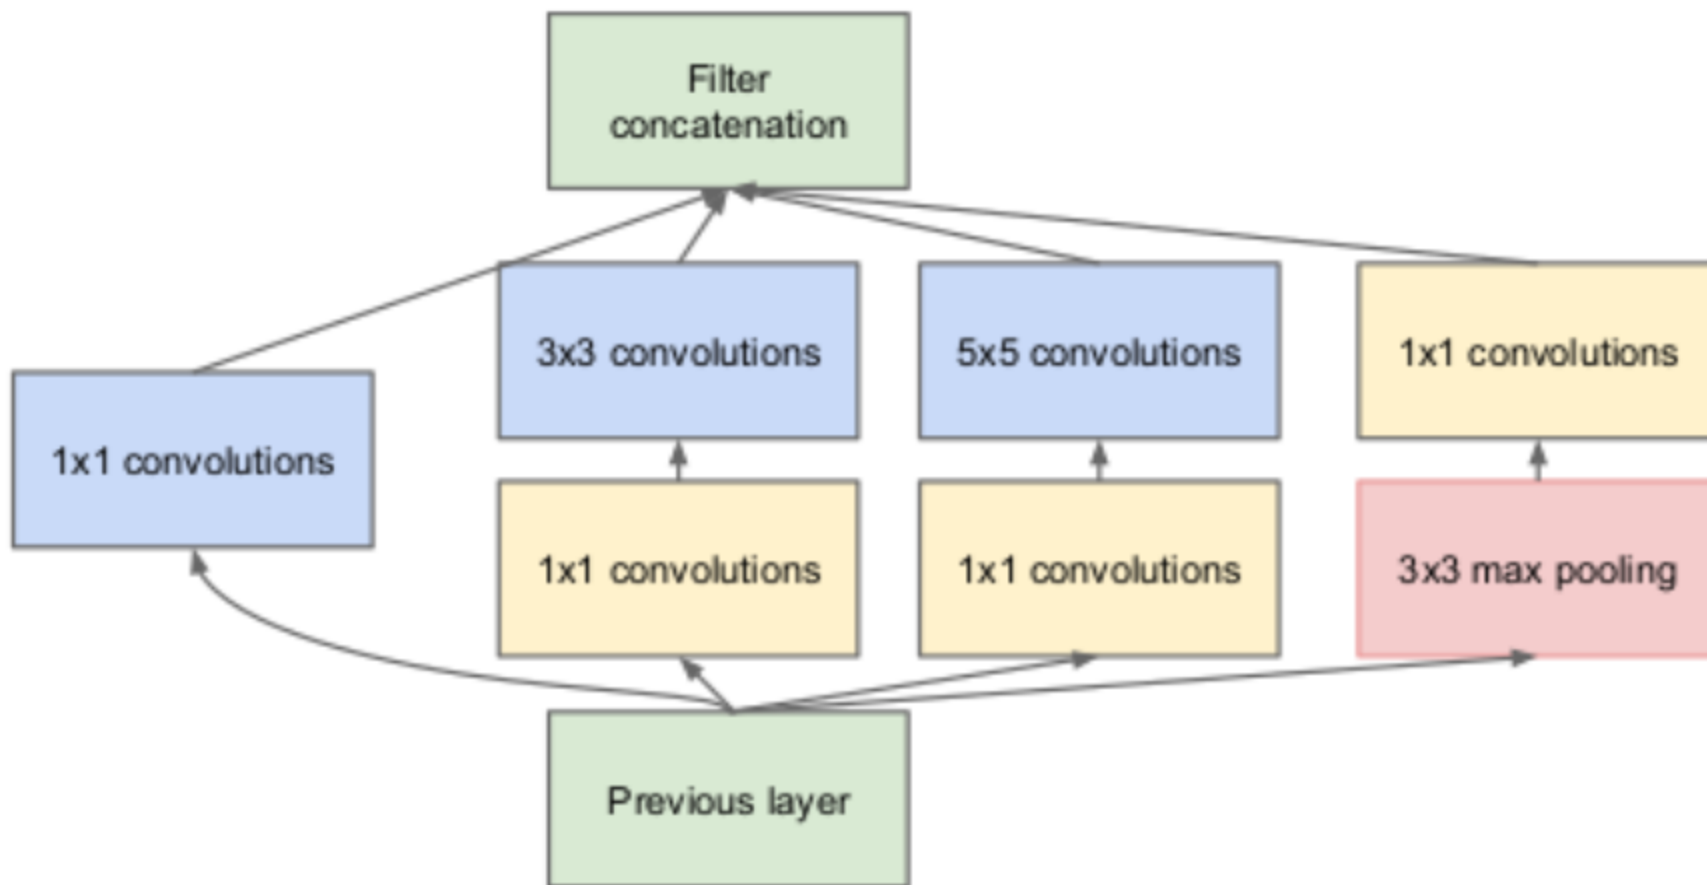


# RESNETS-50

[https://github.com/keras-team/keras-applications/blob/master/keras\\_applications/resnet50.py](https://github.com/keras-team/keras-applications/blob/master/keras_applications/resnet50.py)




# INCEPTION MODULE MOTIVATION



Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., & Wojna, Z. (2016). Rethinking the inception architecture for computer vision. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 2818-2826).

# INCEPTION MODULE MOTIVATION



Input:  $28 \times 28 \times 192$

Filter: Conv  $5 \times 5 \times 192$ , same, 32


Output:  $28 \times 28 \times 32$

Total number of calculations =  $(28 * 28 * 32) * (5 * 5 * 192) = 120 \text{ Million !!}$

# INCEPTION MODULE MOTIVATION


Using 1x1 Convolution to reduce computation cost

A 1x1 convolution is added before the 5x5 convolution  $\Rightarrow$  Also called a **bottleneck layer**



Total number of calculations =  $[(28 * 28 * 16) * (1 * 1 * 192)] + [(28 * 28 * 32) * (5 * 5 * 16)]$   
 $= 12.4 \text{ Million !! (earlier the cost was 120 Million)}$

# GOOGLENET



**Convolution**  
**Pooling**  
**Softmax**  
**Concat/Normalize**

# LeNet-5, AlexNet, VGG-19, GoogLeNet for MNIST Dataset

<http://euler.stat.yale.edu/~tba3/stat665/lectures/lec18/notebook18.html>

# TRANSFER LEARNING

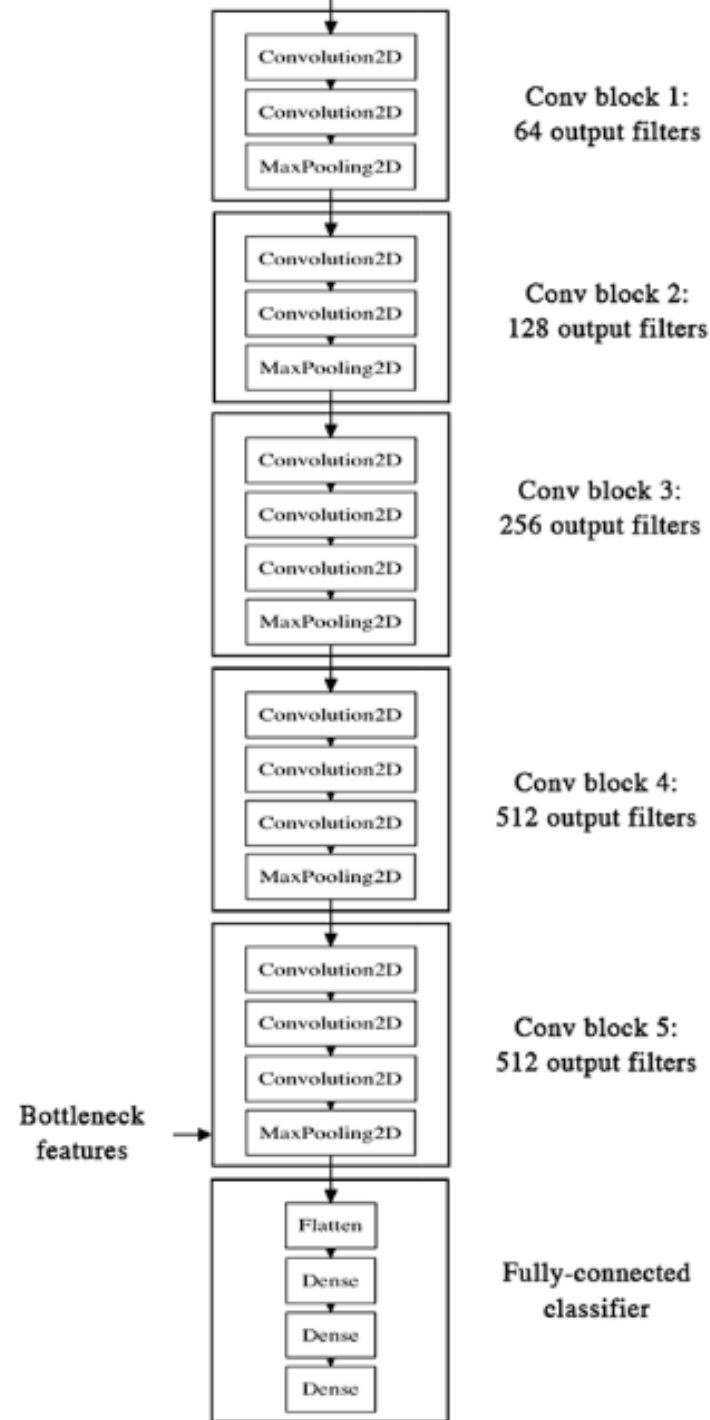
**Transfer learning (TL) is a research problem in machine learning (ML) that focuses on storing knowledge gained while solving one problem and applying it to a different but related problem.**

-Wikipedia

# TRANSFER LEARNING

<https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>

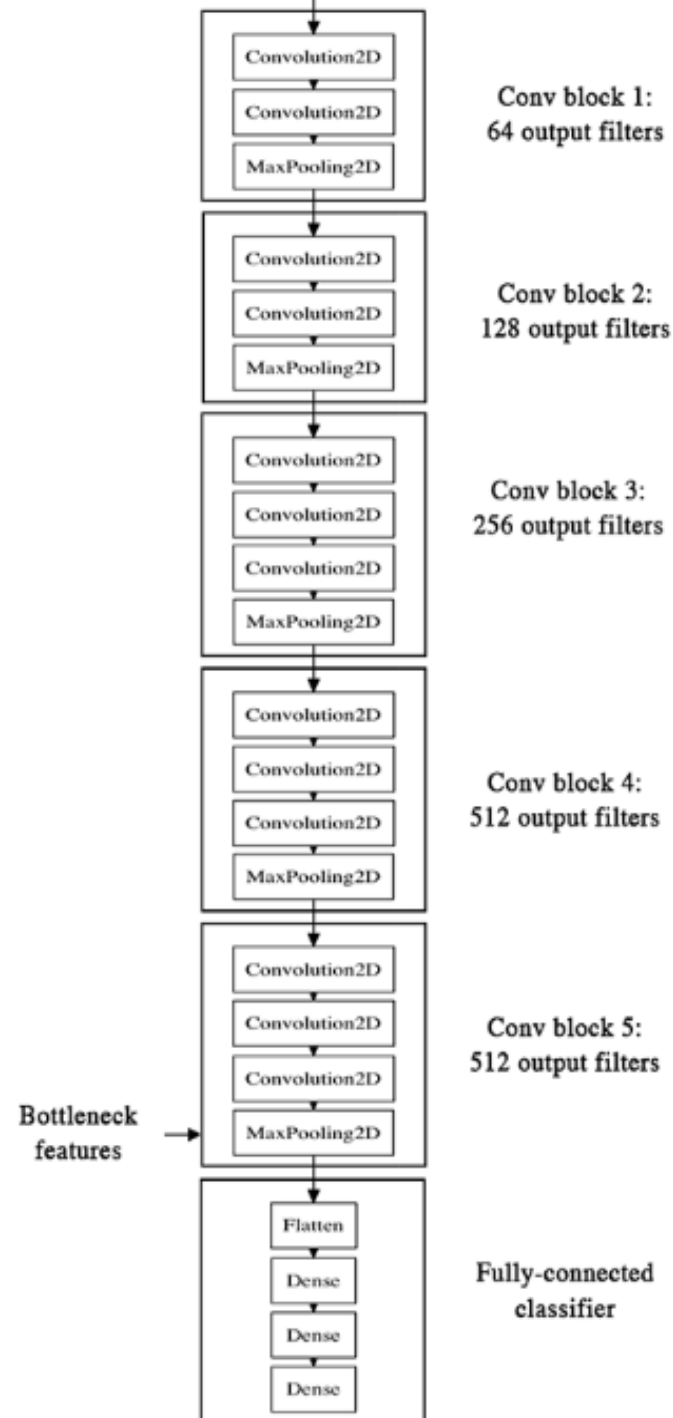
- VGGNet Trained using IMAGENET





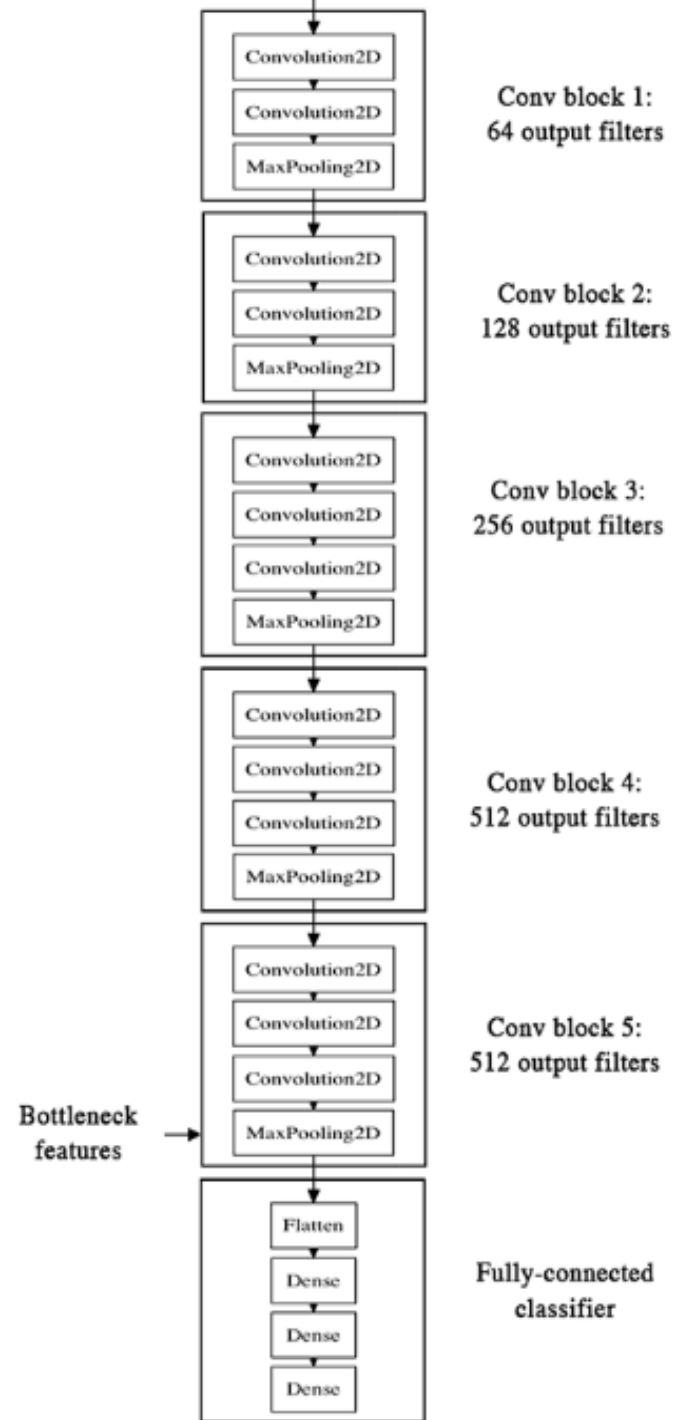
# TRANSFER LEARNING(CASE-0)

- **Pre-Trained CNN Models**
  - **Use the Pre-trained CNN to predict the new dataset**



# TRANSFER LEARNING(CASE-1)


- **Pre-Trained CNN Models + ML Classifiers**
  - **Remove the last Dense Connected Layer**
  - **Take Bottleneck features and use it on a Shallow ML Classifiers**
  - **Here Pre-Trained CNN is used for Feature Engineering**



# TRANSFER LEARNING (CASE-2)

- **Fine Tuning the last two layers of CNN**
  - **Take the Original Dataset and Pretrained CNN Models**
  - **Freeze the early layers (Don't Change)**
  - **Fine tune the last two layers using the new dataset**

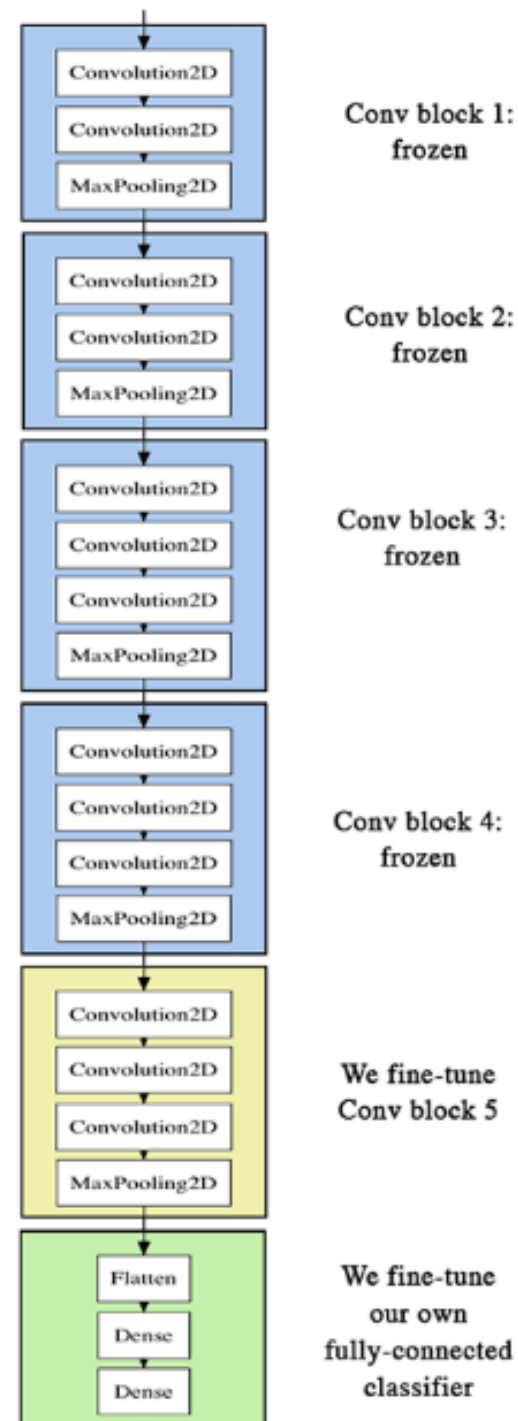
**While fine tuning ensure small learning rate**



# TRANSFER LEARNING (CASE-3)

- **Fine Tuning the Complete model taking the pretrained model as initial model**
  - **Take the Original Dataset and Pretrained CNN Models**
  - **Fine tune the complete model using the new dataset**

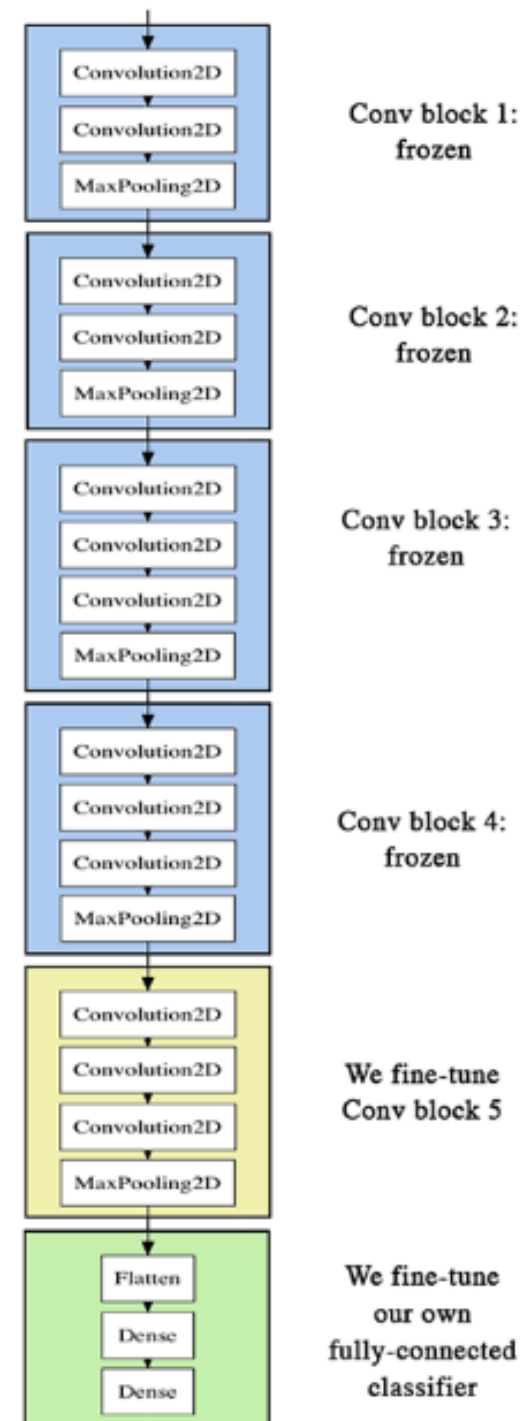
**While fine tuning ensure small learning rate**



# TRANSFER LEARNING (CASE-4)

- Dump everything & Retrain from Scratch

Not widely used



# HOW TO CHOOSE THE TYPE OF TRANSFER LEARNING

- **Based on:**
  - Size of the Dataset
  - Characteristic of the new dataset to the Imagenet Dataset

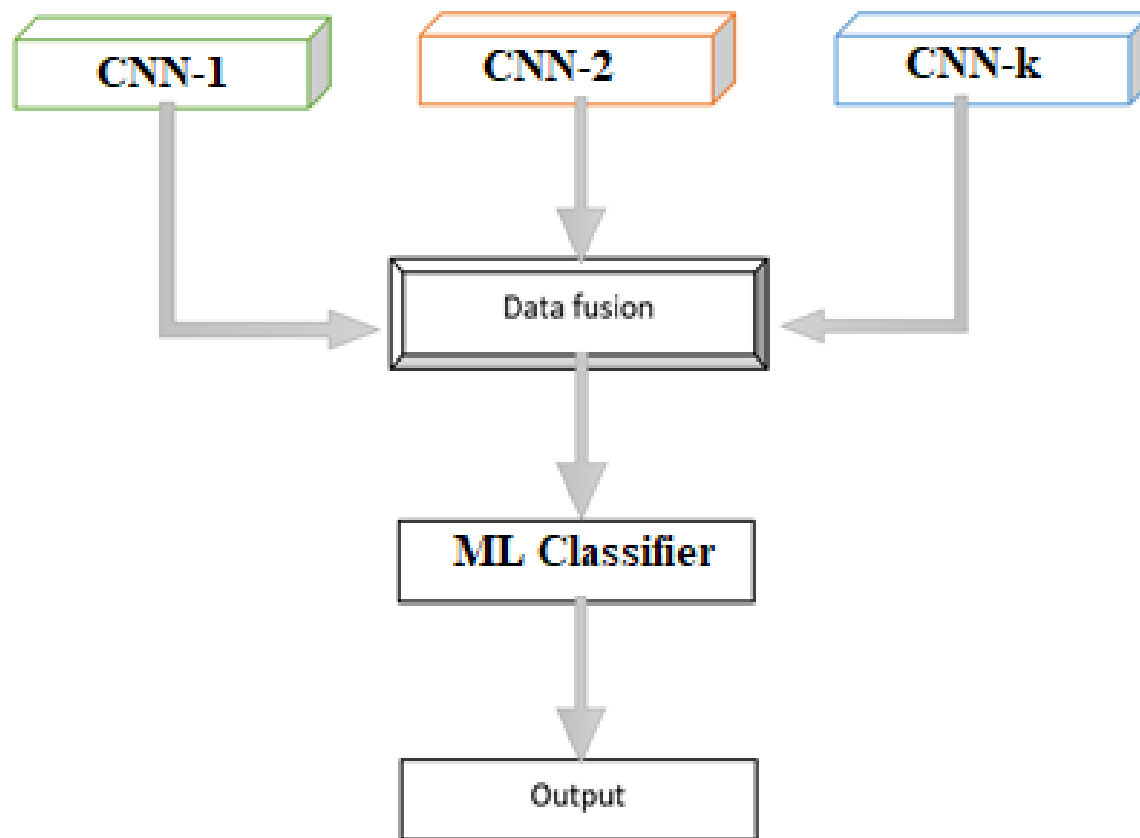
<https://cs231n.github.io/transfer-learning/>

# HOW TO CHOOSE THE TYPE OF TRANSFER LEARNING

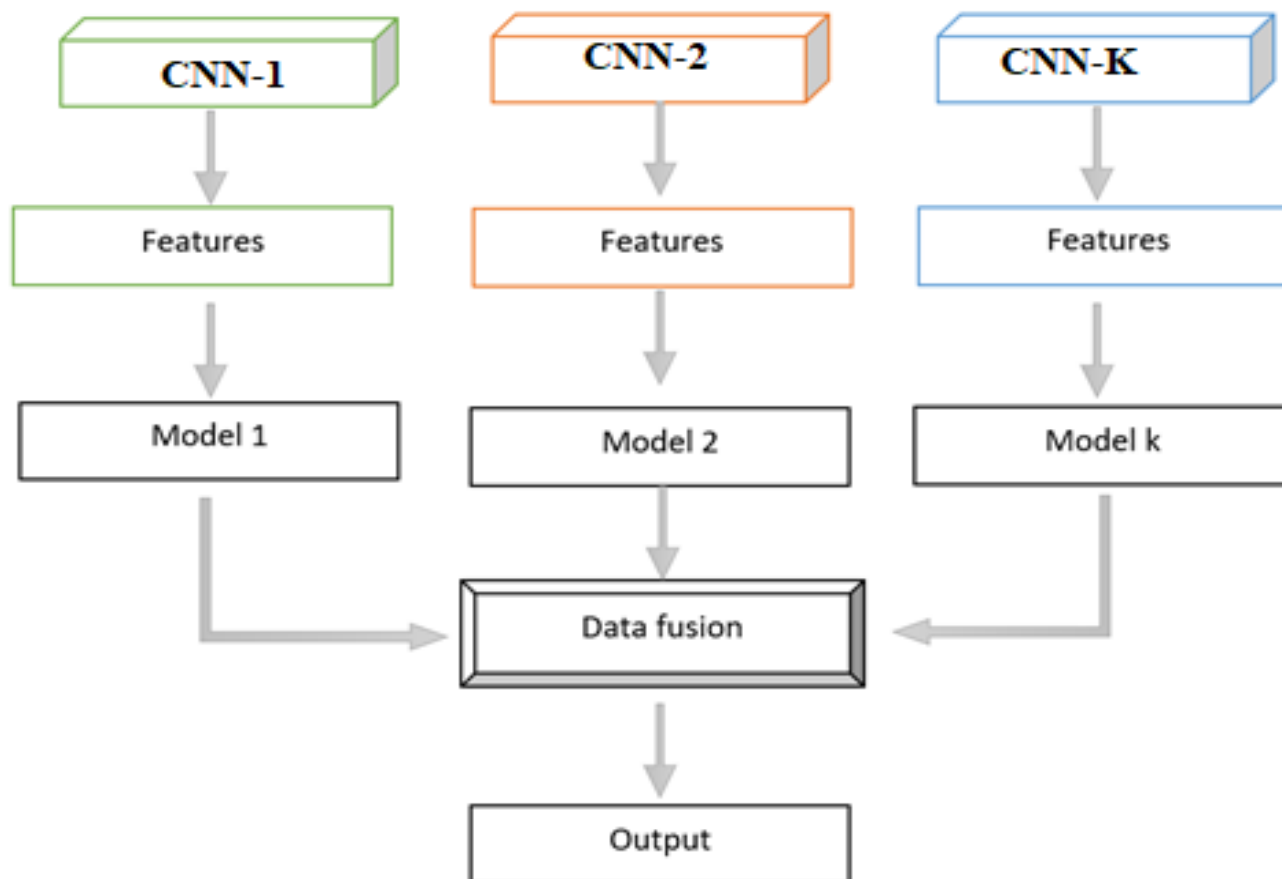
- Case:1
  - If Size(Dataset)=Small and Similar to (IMAGENET)
    - Use Case-1 of Transfer Learning
- Case:2
  - If Size(Dataset)=Large and Similar to (IMAGENET)
    - Use Case-3 of Transfer Learning
- Case:3
  - If Size(Dataset)=Medium and Similar to (IMAGENET)
    - Use Case-2 of Transfer Learning
- Case:4
  - If Size(Dataset)=Small and DisSimilar to (IMAGENET)
    - Use Initial Layers and dump middle layers and use flatten and train a Shallow ML model.
- Case:5
  - If Size(Dataset)=Large and DisSimilar to (IMAGENET)
    - Initialize the Model using pretrained CNN and tune the whole model.




# EARLY FUSION (OR FEATURE LEVEL FUSION)



# LATE FUSION (OR DECISION LEVEL FUSION)



# INTERMEDIATE FUSION





For Your Valuable Time.